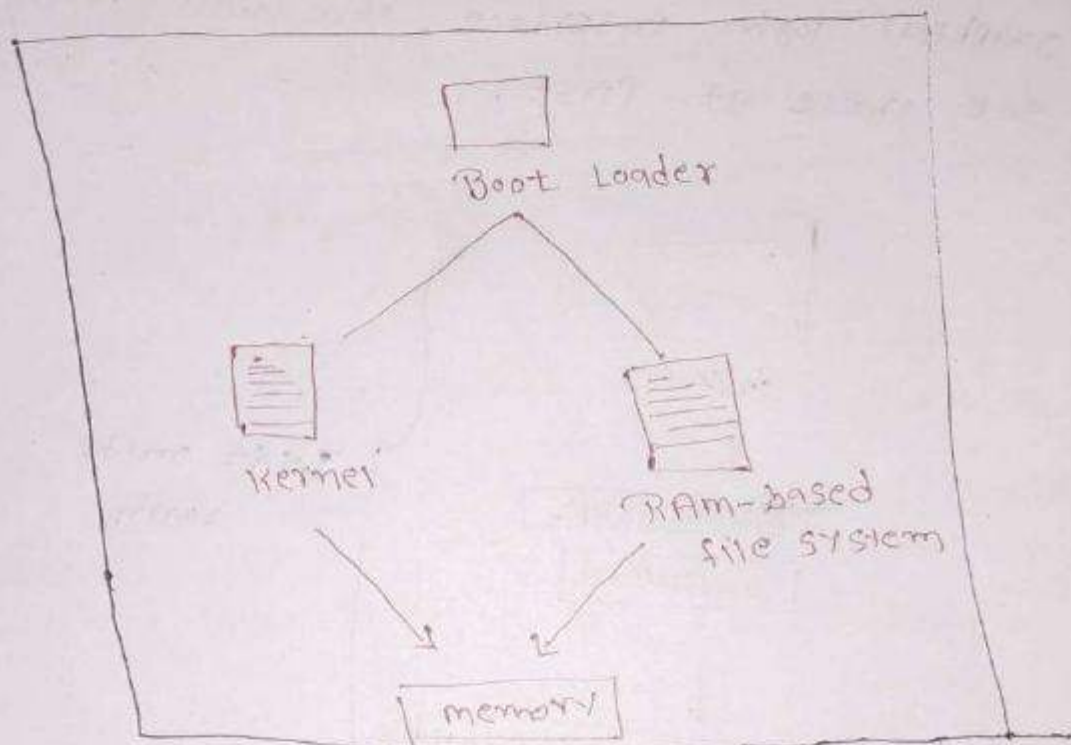


* Kernel, init and services

* The Linux Kernel

→ The boot loads both kernel and an initial RAM-based file system (initramfs) into memory, so it can be used directly by the kernel.



The Linux Kernel

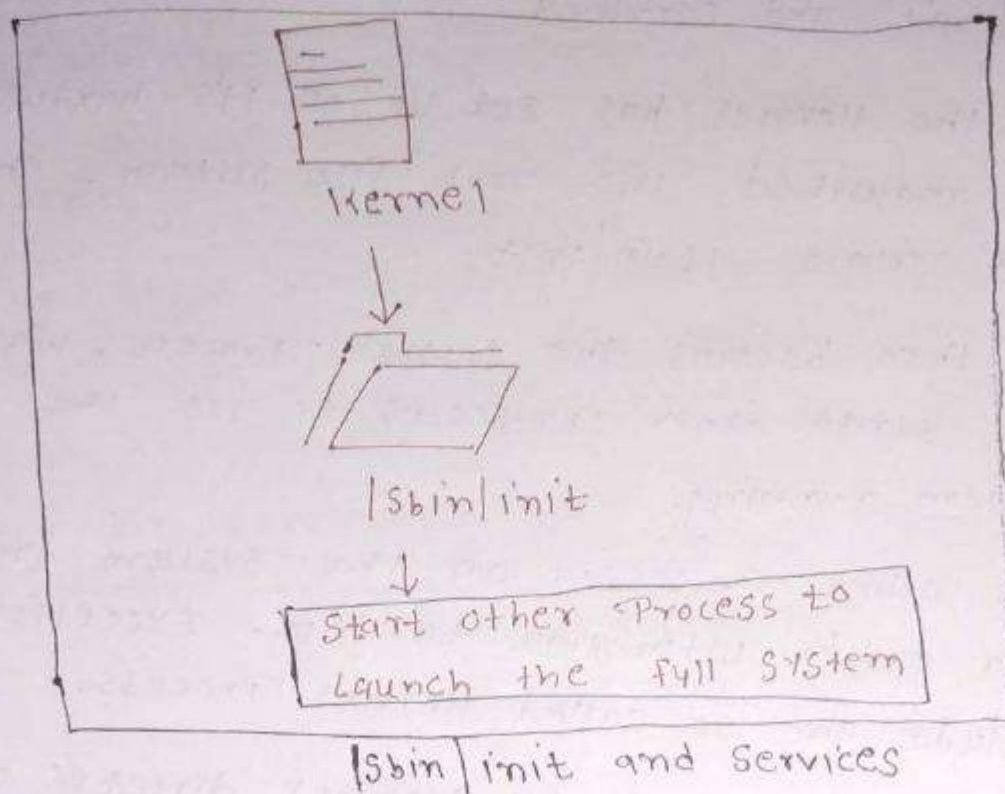
→ when the kernel is loaded in RAM, it immediately initialize and configure the computer's memory and also configures all the hardware attached to the system.

→ The include all processors, I/O subsystem, storage devices etc.

→ The kernel also loads some necessary user space application.

* /sbin/init and services

- once the kernel has set up all its hardware and mounted the root file system, the kernel runs /sbin/init.
- This then becomes the initial process, which then starts other processes to get the system running.
- most other processes on the system trace their origin ultimately to init. Exception include the so called kernel process.
- These are started by kernel directly and their job is to manage internal operating system details.
- Besides starting the system, init is responsible for keeping the system running and for shutting it down cleanly.
- one of its responsibilities is to act when necessary as a manager for all known kernel process
- It cleans up after them upon completion and restarts user log in services as needed when user log in and out and does the same for other background system services.



- Traditionally, this process startup was done using convention that date back to the 1980's and the system v variety of UNIX.
- This Serial Process (called sysvinit) had the system pass through a sequence of runlevels containing collection of scripts that start and stop services.
- Each runlevel supported a different mode of running the system.
- Within each runlevel, individual could be set to run, or to be shut down if running.
- However, all major distribution have moved away from this sequential method of system initialization, although they usually can emulate many system v utilities for compatibility purpose.

* StartUP Alternatives

- Sysvinit viewed things as serial process, divided into a series of sequential stages.
- Each stage required completion before next could proceed.
- Thus, start up did not easily take advantage of parallel processing that could be done with multiple processors or cores found on modern systems.
- Furthermore, starting up and rebooting were seen as relatively rare events; exactly how long they took was not considered important.
- This is no longer true, especially with mobile devices are embedded Linux systems.
- Some modern methods, such as the use of containers, can require almost instantaneous start up time.
- Thus systems now require methods with faster and enhanced capabilities.
- Finally the older methods require rather complicated start up scripts
- which were difficult to keep universal across distribution version, kernel versions, architectures, and the types of systems.
- The two main alternatives developed were:

UPstart

- Developed by Ubuntu and first included in 2006
- Adopted in Fedora 9 (in 2008) and in RHEL 6 and its clones

Systemd

- Adopted by Fedora first (in 2011)
- Adopted by RHEL 7 and SUSE
- Replace UPstart in Ubuntu 16.04

→ while the migration to systemd was rather controversial, it has been adopted by all major distributions,

so

→ Regardless of how one feels about the controversies of the technical method of systemd, almost universal adoption has made learning how to work on Linux systems simpler

* Systemd Features

- Systems with systemd start up faster than those with earlier init methods.
- This is largely because it replaced a serialized set of steps with aggressive Parallelization techniques
- With permits multiple services to be initiated simultaneously.
- Complicated startup shell scripts are replaced with simpler configuration files, which enumerate what has to be done before a service ~~start~~ started.
- How to executor service startup.
- And what condition the service should indicate have been accomplished when startup is finished.
- one thing to note is that `/sbin/init` now just points to `/lib/systemd/systemd`,
i.e. systemd takes over the init process.
- one systemd command (`systemctl`) is used for most basic task.

- starting, stoping, restarting a service (using httpd, the Apache web server as an example) on currently running system,

```
$ sudo systemctl  
start | stop | restart  
httpd.service
```

- Enabling or disabling a system service from starting up at system boot

```
$ sudo systemctl enable | disable  
httpd.service
```

- checking on the status of a service

```
$ sudo systemctl status  
httpd.service
```

→ In most cases, the service can be omitted.

→ There are many technical difference with older methods that lie beyond the scope of our discussion.

Systemd

* Lab :- Apache we Server Status

- check the status of httpd (Apache web server) on your system
- If it is running, stop it and check again.
- start the service and check the status
- you probably want to stop it again when you are done.

1.	\$	sudo	systemctl	status	httpd
2.	\$	sudo	systemctl	start	httpd
3.	\$	sudo	systemctl	status	httpd
4.	\$	sudo	systemctl	stop	httpd

↑
Solution

* Linux File Systems

- Libraries separate books and other media into multiple sections; this organization will depend on the subject matter, audience, media type, and frequency of retrieval.
- The same concept applies to a file system which is the embodiment of method of storing and organizing arbitrary collections of data in a human-usable form.
- Different types of file systems supported by Linux:
 - Conventional disk file systems
 - ext3, ext4, XFS, Btrfs, JFS, NTFS, vfat, exfat, etc.
 - Flash storage file systems
 - ubifs, jffs2, yaffs, etc.
 - Special purpose file systems
 - Procfs, sysfs, tmpfs, squashfs, debugfs, fuse, etc.

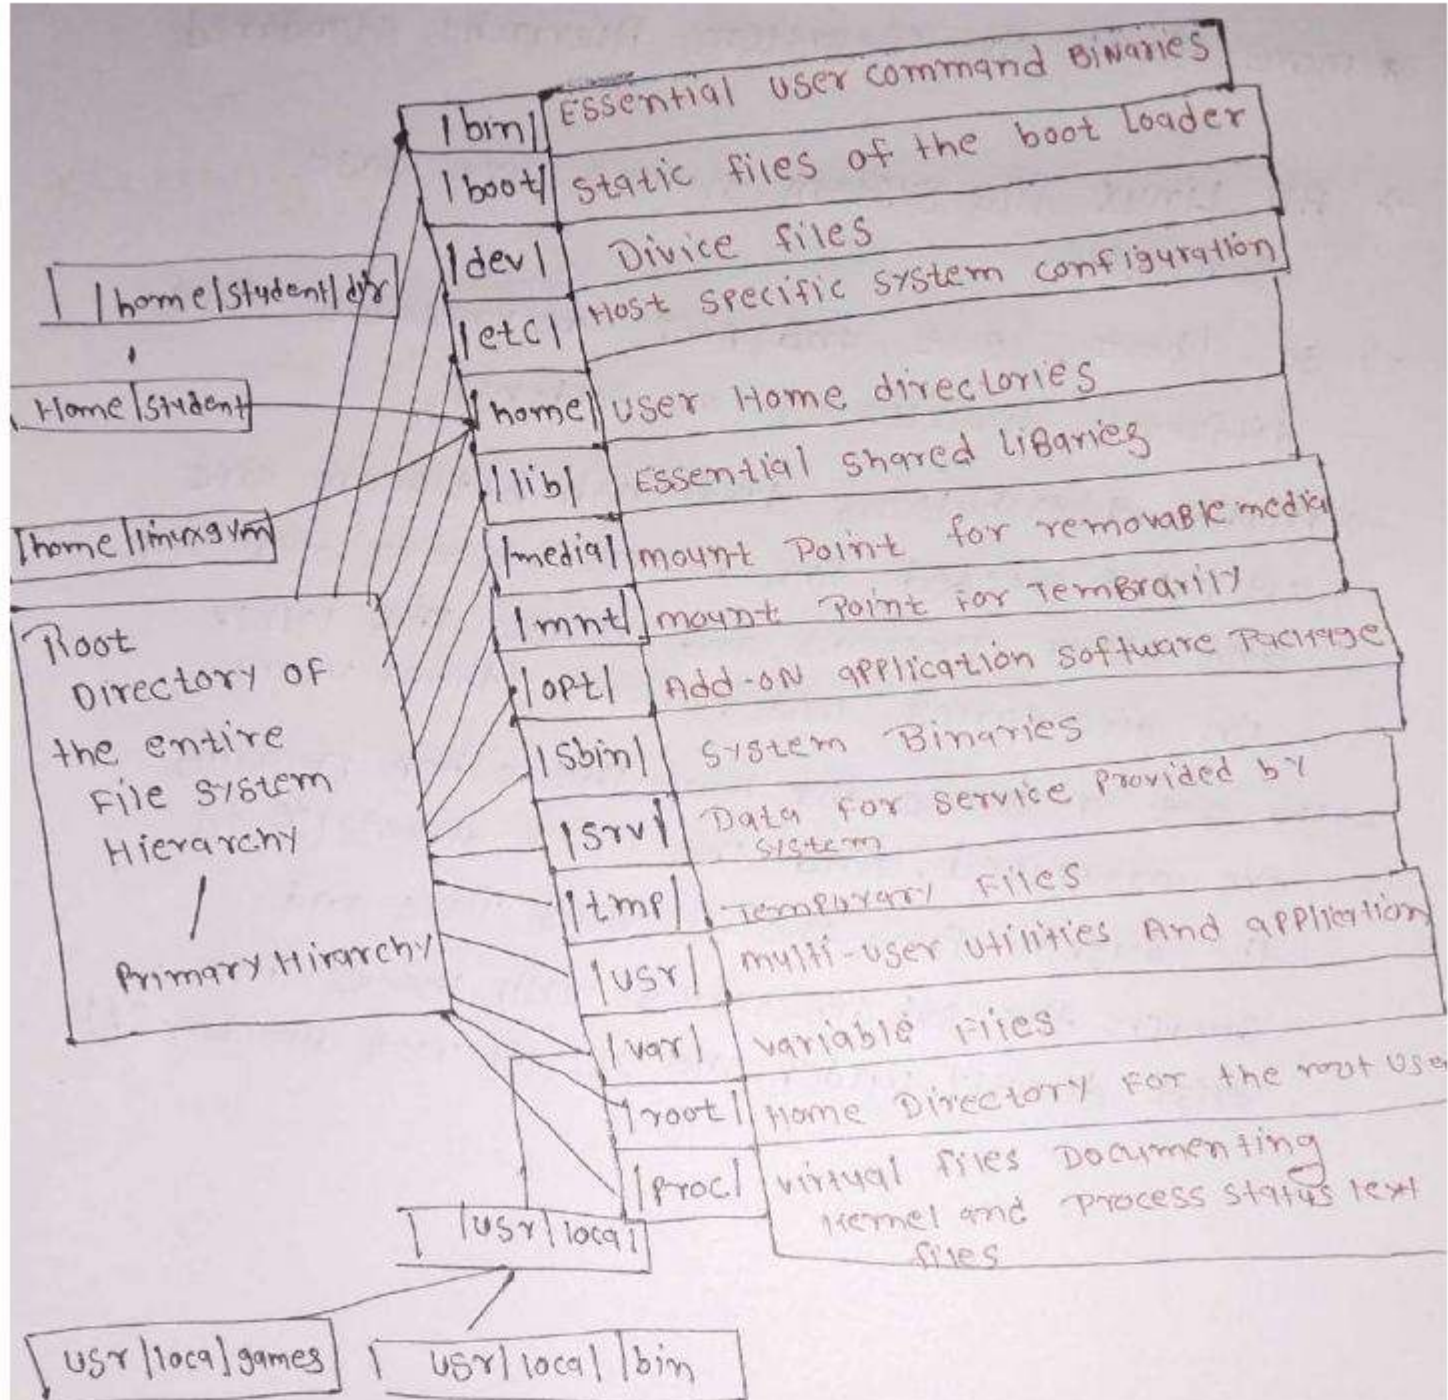
* Partitions and File systems

- A partition is a dedicated subsection of physical storage.
- Historically this meant a physically contiguous portion of hard disk
- Today's storage devices can be more complicated but we still think of a partition as a fixed area to be treated as a whole.
- A file system is just a method of storing and accessing files.
- one can think of a partition as a container in which a file system resides.
- However in some circumstances, a file system can span more than one partition if one uses symbolic links,

	windows	Linux
Partition	DISK1	/dev/sda1
Filesystem Type	NTFS/VFAT	EXT3/EXT4/XFS/BTRFS
mounting Parameters	Drive Letter	mount Point
Base folder	c:\	/

* The Filesystem Hierarchy Standard

- Linux systems store their important files according to standard layout called the Filesystem Hierarchy Standard (FHS).
- which has long been maintained by the Linux Foundation.
- Linux uses '/' character to separate paths (as is UNIX unlike windows, which uses '\') and does not have drive letters.
- multiple drives and/or Partitions are mounted as directories in the single file system.
- Removable media such as USB drives and CDs and DVDs will show up as mounted at
`/run/media/yourusername/disklabel` for recent Linux systems or under `/media` for older distributions.
- For example, if your username is student, a USB Pen drive labeled FEDORA might end up being found at
`/run/media/student/FEDORA`, and a file `README.txt` on that disc would be at
`/run/media/student/FEDORA/README.txt`.



File System Hierarchy Standard (FHS)

* more About the Filesystem Hierarchy Standard

→ All Linux file system names are case-sensitive.

→ So `/boot`, `/Boot`, and `/BOOT` represent three different directories (or folders).

→ many distributions distinguish between core utilities needed for proper system operation and other programs, and place the latter in directories under `/usr` (thru `/usr`).

→ To get a sense for how the other programs are organized, find the `/usr` directory in the diagram from the previous page and compare the subdirectories with those that exist directly under the system root directory (`/`).